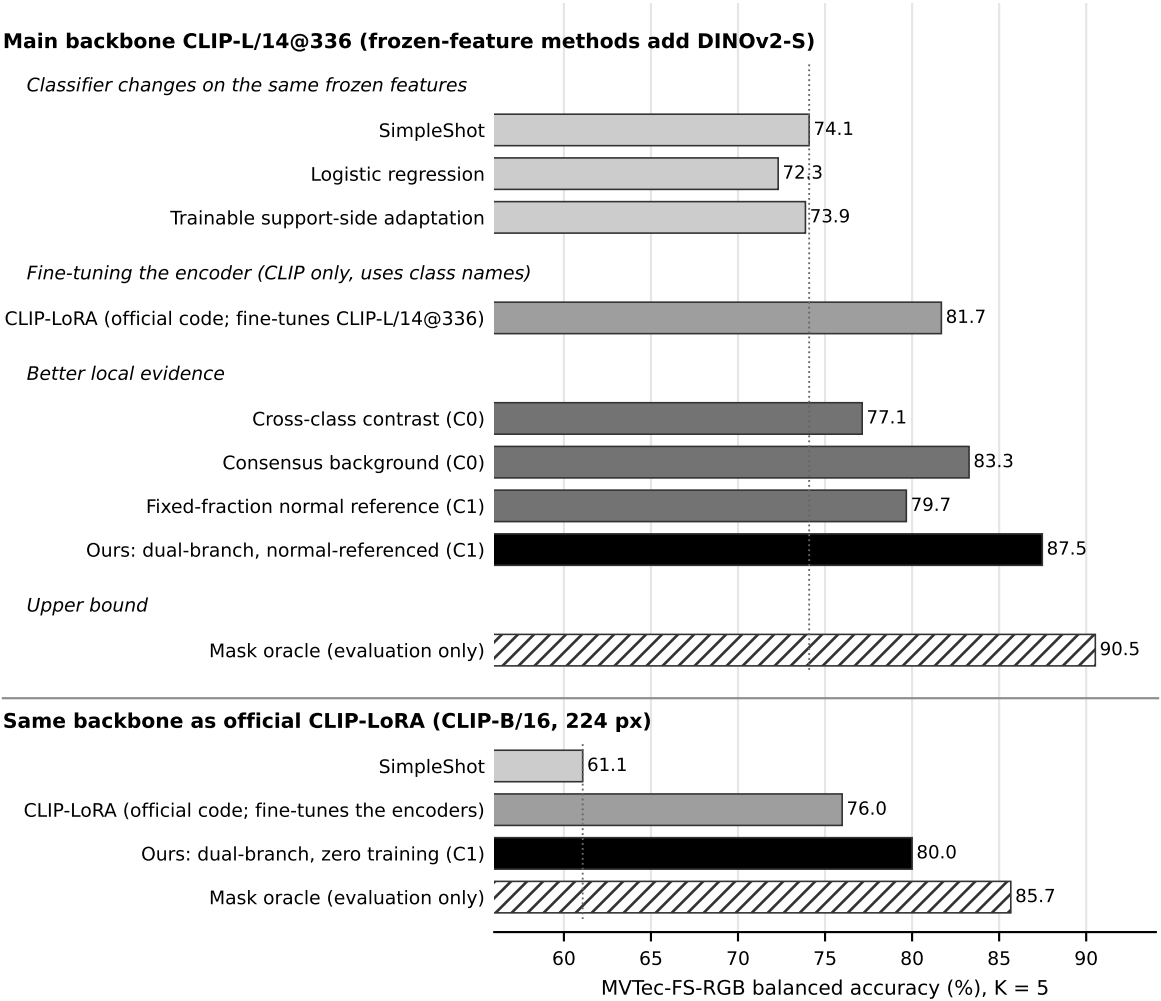
Main backbone CLIP-L/14@336 (frozen-feature methods add DINOv2-S)
Classifier changes on the same frozen features
SimpleShot
74.1
Logistic regression
72.3
Trainable support-side adaptation
73.9
Fine-tuning the encoder (CLIP only, uses class names)
CLIP-LoRA (official code; fine-tunes CLIP-L/14@336)
81.7
Better local evidence
Cross-class contrast (C0)
77.1
Consensus background (C0)
83.3
Fixed-fraction normal reference (C1)
79.7
Ours: dual-branch, normal-referenced (C1)
87.5
Upper bound
Mask oracle (evaluation only)
90.5
Same backbone as official CLIP-LoRA (CLIP-B/16, 224 px)
SimpleShot
61.1
CLIP-LoRA (official code; fine-tunes the encoders)
76.0
Ours: dual-branch, zero training (C1)
80.0
Mask oracle (evaluation only)
85.7
60
65
70
75
80
85
90
MVTec-FS-RGB balanced accuracy (%), K = 5