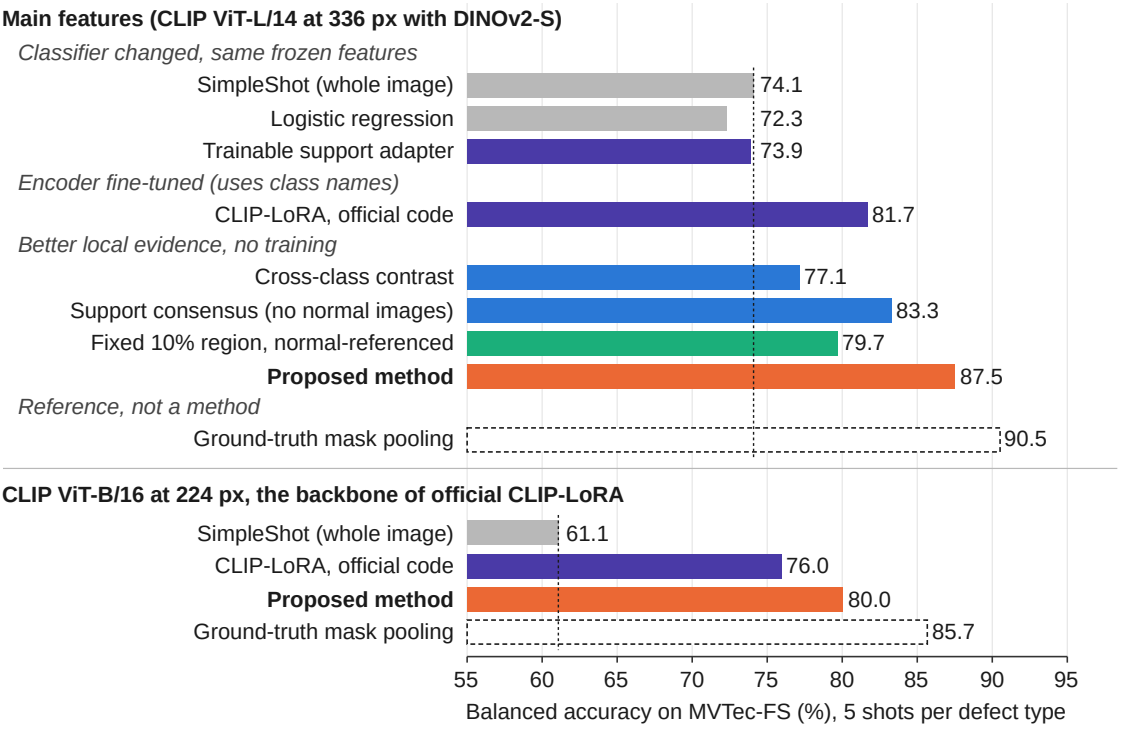
Main features (CLIP ViT-L/14 at 336 px with DINOv2-S)
Classifier changed, same frozen features
SimpleShot (whole image)
74.1
Logistic regression
72.3
Trainable support adapter
73.9
Encoder fine-tuned (uses class names)
CLIP-LoRA, official code
81.7
Better local evidence, no training
Cross-class contrast
77.1
Support consensus (no normal images)
83.3
Fixed 10% region, normal-referenced
79.7
Proposed method
87.5
Reference, not a method
Ground-truth mask pooling
90.5
CLIP ViT-B/16 at 224 px, the backbone of official CLIP-LoRA
SimpleShot (whole image)
61.1
CLIP-LoRA, official code
76.0
Proposed method
80.0
Ground-truth mask pooling
85.7
55
60
65
70
75
80
85
90
95
Balanced accuracy on MVTec-FS (%), 5 shots per defect type